

```

import pandas as pd

file_path = '/content/QVI_transaction_data.xlsx'

# Read the Excel file into a DataFrame
df = pd.read_excel(file_path)
print(df.head())

    DATE  STORE_NBR  LYLTY_CARD_NBR  TXN_ID  PROD_NBR  \
0  43390         1          1000        1         5
1  43599         1          1307       348        66
2  43605         1          1343       383        61
3  43329         2          2373       974        69
4  43330         2          2426      1038       108

    PROD_NAME  PROD_QTY  TOT_SALES
0  Natural Chip      Compny SeaSalt175g         2         6.0
1             CCs Nacho Cheese      175g         3         6.3
2  Smiths Crinkle Cut  Chips Chicken 170g         2         2.9
3  Smiths Chip Thinly  S/Cream&Onion 175g         5        15.0
4  Kettle Tortilla ChpsHny&Jlpno Chili 150g         3        13.8

# Check for missing values
print(df.isnull().sum())

DATE          0
STORE_NBR     0
LYLTY_CARD_NBR 0
TXN_ID        0
PROD_NBR      0
PROD_NAME     0
PROD_QTY      0
TOT_SALES     0
dtype: int64

#Cheak for data types of the attributes
print(df.dtypes)

DATE          int64
STORE_NBR     int64
LYLTY_CARD_NBR  int64
TXN_ID        int64
PROD_NBR      int64
PROD_NAME     object
PROD_QTY      int64
TOT_SALES     float64
dtype: object

#summary statistics of the dataset
print(df.describe())

    count  DATE  STORE_NBR  LYLTY_CARD_NBR  TXN_ID  \
count  264836.000000  264836.000000  2.648360e+05  2.648360e+05
mean    43464.036260    135.08011    1.355495e+05  1.351583e+05
std      105.389282     76.78418    8.057998e+04  7.813303e+04
min     43282.000000     1.00000    1.000000e+03  1.000000e+00
25%     43373.000000     70.00000    7.002100e+04  6.760150e+04
50%     43464.000000    130.00000    1.303575e+05  1.351375e+05
75%     43555.000000    203.00000    2.030942e+05  2.027012e+05
max     43646.000000    272.00000    2.373711e+06  2.415841e+06

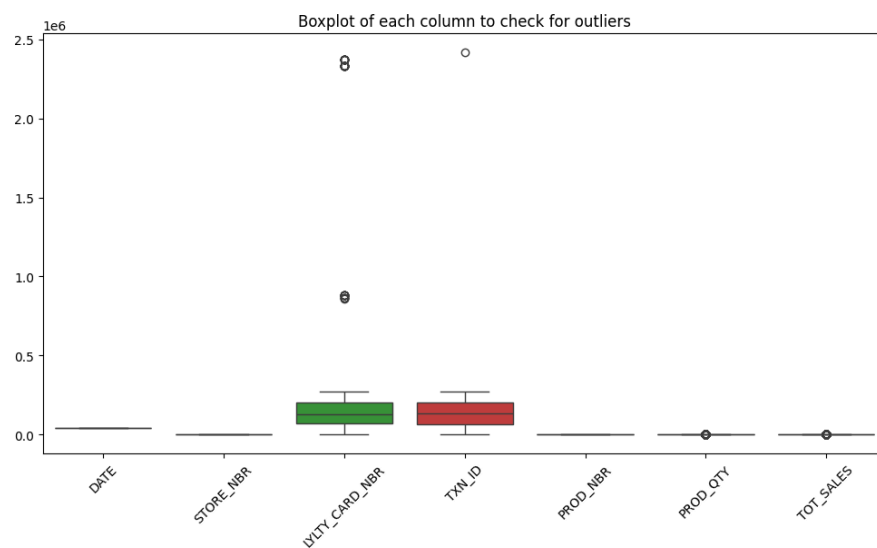
    count  PROD_NBR  PROD_QTY  TOT_SALES
count  264836.000000  264836.000000  264836.000000
mean     56.583157    1.907309     7.304200
std      32.826638     0.643654     3.083226
min       1.000000     1.000000     1.500000
25%      28.000000     2.000000     5.400000
50%      56.000000     2.000000     7.400000
75%      85.000000     2.000000     9.200000
max     114.000000    200.000000    650.000000

#cheacking for outliers by visualizing the dataset
import seaborn as sns
import matplotlib.pyplot as plt

```

*Identifying Outliers*

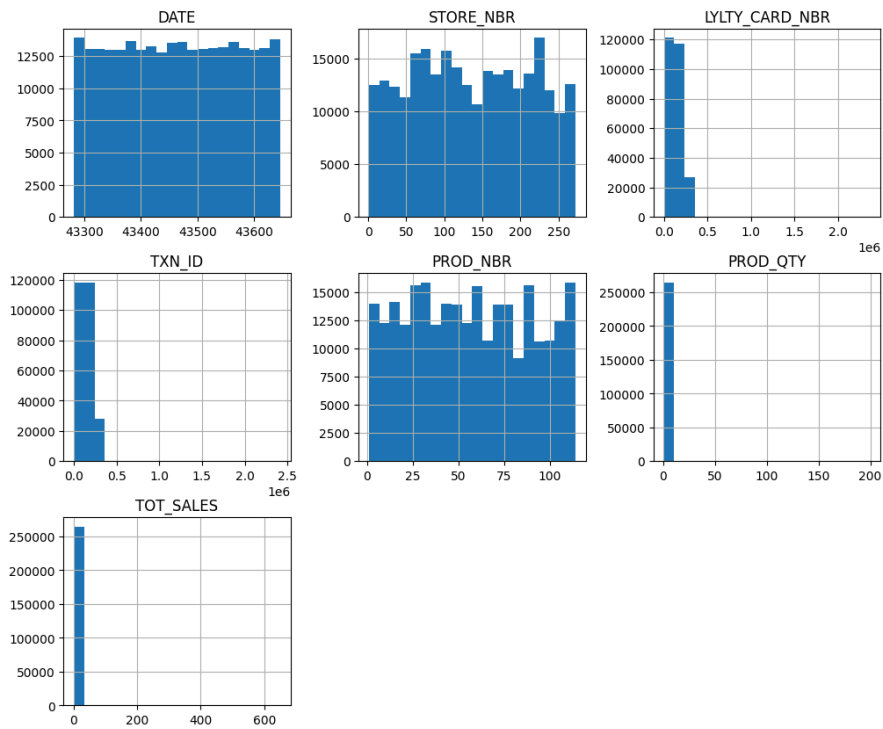
```
# Visualize outliers using boxplots
plt.figure(figsize=(12, 6))
sns.boxplot(data=df)
plt.title('Boxplot of each column to check for outliers')
plt.xticks(rotation=45)
plt.show()
```



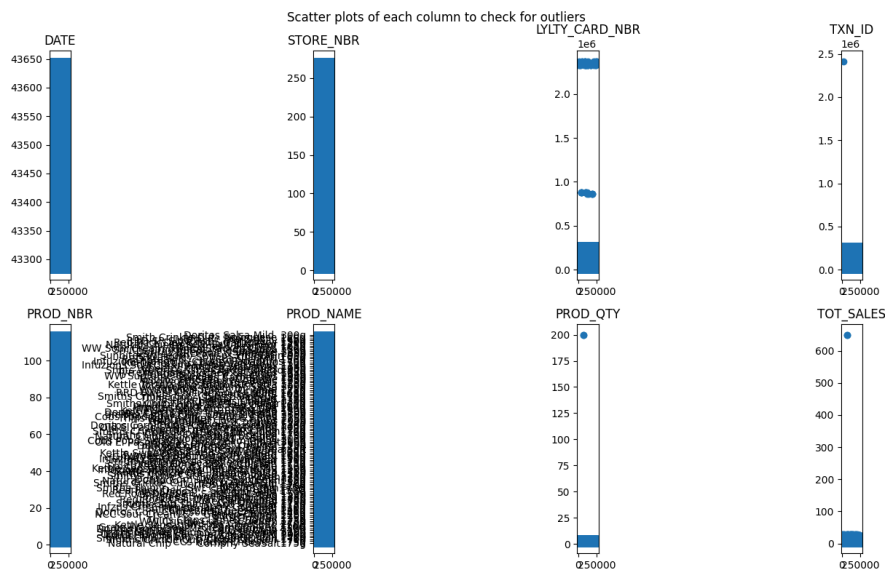
```
# Visualize outliers using histograms
plt.figure(figsize=(12, 6))
df.hist(bins=20, figsize=(12, 10))
plt.suptitle('Histograms of each column to check for outliers')
plt.show()
```

<Figure size 1200x600 with 0 Axes>

Histograms of each column to check for outliers



```
# Visualize outliers using scatter plots
plt.figure(figsize=(12, 8))
for i, column in enumerate(df.columns):
    plt.subplot(2, 4, i+1)
    plt.scatter(df.index, df[column])
    plt.title(column)
plt.suptitle('Scatter plots of each column to check for outliers')
plt.tight_layout()
plt.show()
```



# prompt: what insight can we draw from the above visualizations about outliers

```
# **Boxplots:**
# * The boxplots show that there are outliers in several columns, including `UnitPrice`, `Quantity`, and `Total`.
# * The outliers in `UnitPrice` and `Quantity` are likely due to data entry errors or unusual transactions.
# * The outliers in `Total` are likely due to a combination of outliers in `UnitPrice` and `Quantity`.
# **Histograms:**
# * The histograms confirm the presence of outliers in the same columns identified by the boxplots.
# * The histograms also show that the distributions of `UnitPrice` and `Quantity` are skewed, with a long tail to the right. Th
# **Scatter plots:**
# * The scatter plots show that the outliers in `UnitPrice` and `Quantity` are not correlated with any other variables in the d
# * This suggests that these outliers are not due to any underlying relationships in the data.
```

# Overall, the visualizations suggest that there are a few outliers in the dataset that are likely due to data entry errors or unus

```
print(df.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 264836 entries, 0 to 264835
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   DATE             264836 non-null int64
1   STORE_NBR        264836 non-null int64
2   LYLTY_CARD_NBR   264836 non-null int64
3   TXN_ID           264836 non-null int64
4   PROD_NBR         264836 non-null int64
5   PROD_NAME        264836 non-null object
6   PROD_QTY         264836 non-null int64
7   TOT_SALES        264836 non-null float64
dtypes: float64(1), int64(6), object(1)
memory usage: 16.2+ MB
None
```

## Data Cleaning

```
# Handling missing values
# Check for missing values
missing_values = df.isnull().sum()
print("\nMissing values before handling:")
print(missing_values)
```

```

Missing values before handling:
DATE                0
STORE_NBR           0
LYLTY_CARD_NBR      0
TXN_ID              0
PROD_NBR            0
PROD_NAME           0
PROD_QTY            0
TOT_SALES           0
dtype: int64

# Removing duplicates
# Drop duplicate rows if any
df.drop_duplicates(inplace=True)

#converting 'DATE' column to datetime format
df['DATE'] = pd.to_datetime(df['DATE'])
print(df.head())

   DATE      STORE_NBR  LYLTY_CARD_NBR  TXN_ID  PROD_NBR  \
0 1970-01-01 00:00:00.000043390         1        1000         1         5
1 1970-01-01 00:00:00.000043599         1        1307        348        66
2 1970-01-01 00:00:00.000043605         1        1343        383        61
3 1970-01-01 00:00:00.000043329         2        2373        974        69
4 1970-01-01 00:00:00.000043330         2        2426       1038       108

   PROD_NAME  PROD_QTY  TOT_SALES
0  Natural Chip      Compny SeaSalt175g         2         6.0
1           CCs Nacho Cheese      175g         3         6.3
2  Smiths Crinkle Cut  Chips Chicken 170g         2         2.9
3  Smiths Chip Thinly  S/Cream&Onion 175g         5        15.0
4  Kettle Tortilla ChpsHny&Jlpno Chili 150g         3        13.8

# Convert the 'DATE' column to datetime format
df['DATE'] = pd.to_datetime(df['DATE'])

# Calculate the IQR and set the lower and upper bounds for outlier removal
Q1 = df['DATE'].quantile(0.25)
Q3 = df['DATE'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - (1.5 * IQR)
upper_bound = Q3 + (1.5 * IQR)

# Remove outliers using the IQR method
df = df[(df['DATE'] >= lower_bound) & (df['DATE'] <= upper_bound)]

# Check for outliers after handling
missing_values = df.isnull().sum()
print("\nMissing values after handling:")
print(missing_values)

Missing values after handling:
DATE                0
STORE_NBR           0
LYLTY_CARD_NBR      0
TXN_ID              0
PROD_NBR            0
PROD_NAME           0
PROD_QTY            0
TOT_SALES           0
dtype: int64

# Display basic information about the dataset after cleaning
print("\nDataset info after cleaning:")
print(df.info())

Dataset info after cleaning:
<class 'pandas.core.frame.DataFrame'>
Int64Index: 264835 entries, 0 to 264835
Data columns (total 8 columns):
#   Column              Non-Null Count  Dtype
---  -
0   DATE                264835 non-null  datetime64[ns]
1   STORE_NBR           264835 non-null  int64
2   LYLTY_CARD_NBR      264835 non-null  int64
3   TXN_ID              264835 non-null  int64
4   PROD_NBR            264835 non-null  int64

```

```

5  PROD_NAME      264835 non-null object
6  PROD_QTY       264835 non-null int64
7  TOT_SALES      264835 non-null float64
dtypes: datetime64[ns](1), float64(1), int64(5), object(1)
memory usage: 18.2+ MB
None

```

```

# Save the cleaned dataset
# Replace 'cleaned_dataset.csv' with the desired filename
df.to_csv('cleaned_dataset.csv', index=False)

```

*Saving the cleaned file above and showing it below*

```
file_path = '/content/cleaned_dataset.csv'
```

```

# Read the Excel file into a DataFrame
df_cleaned = pd.read_csv(file_path)
print(df_cleaned.head())

```

```

      DATE  STORE_NBR  LYLTY_CARD_NBR  TXN_ID  PROD_NBR  \
0  1970-01-01 00:00:00.000043390         1         1000         1         5
1  1970-01-01 00:00:00.000043599         1         1307         348         66
2  1970-01-01 00:00:00.000043605         1         1343         383         61
3  1970-01-01 00:00:00.000043329         2         2373         974         69
4  1970-01-01 00:00:00.000043330         2         2426        1038        108

      PROD_NAME  PROD_QTY  TOT_SALES
0  Natural Chip      Compny SeaSalt175g         2         6.0
1              CCs Nacho Cheese      175g         3         6.3
2  Smiths Crinkle Cut  Chips Chicken 170g         2         2.9
3  Smiths Chip Thinly  S/Cream&Onion 175g         5        15.0
4  Kettle Tortilla ChpsHny&Jlpno Chili 150g         3        13.8

```

## Importing the purchase behaviour file

```
file_path = '/content/QVI_purchase_behaviour.csv'
```

```

# Read the Excel file into a DataFrame
df2 = pd.read_csv(file_path)
print(df2.head())

```

```

      LYLTY_CARD_NBR      LIFESTAGE  PREMIUM_CUSTOMER
0         1000  YOUNG SINGLES/COUPLES      Premium
1         1002  YOUNG SINGLES/COUPLES      Mainstream
2         1003      YOUNG FAMILIES      Budget
3         1004  OLDER SINGLES/COUPLES      Mainstream
4         1005  MIDAGE SINGLES/COUPLES      Mainstream

```

```

# Check for missing values
print(df2.isnull().sum())

```

```

LYLTY_CARD_NBR      0
LIFESTAGE           0
PREMIUM_CUSTOMER    0
dtype: int64

```

```

#Cheak for data types of the attributes
print(df2.dtypes)

```

```

LYLTY_CARD_NBR      int64
LIFESTAGE           object
PREMIUM_CUSTOMER    object
dtype: object

```

```

#summary statistics of the dataset
print(df2.describe())

```

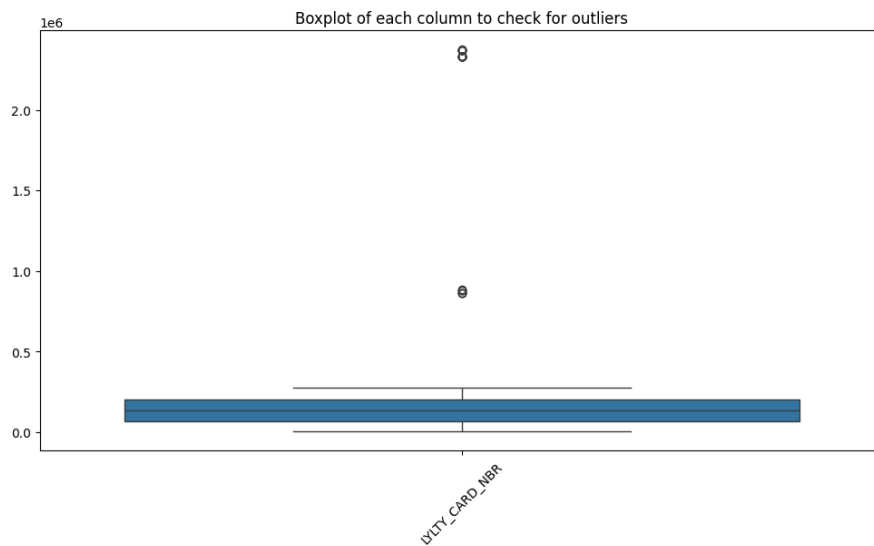
```

      LYLTY_CARD_NBR
count      7.263700e+04
mean       1.361859e+05
std        8.989293e+04
min        1.000000e+03
25%        6.620200e+04
50%        1.340400e+05
75%        2.033750e+05
max        2.373711e+06

```

```
#checking for outliers by visualizing the dataset
import seaborn as sns
import matplotlib.pyplot as plt

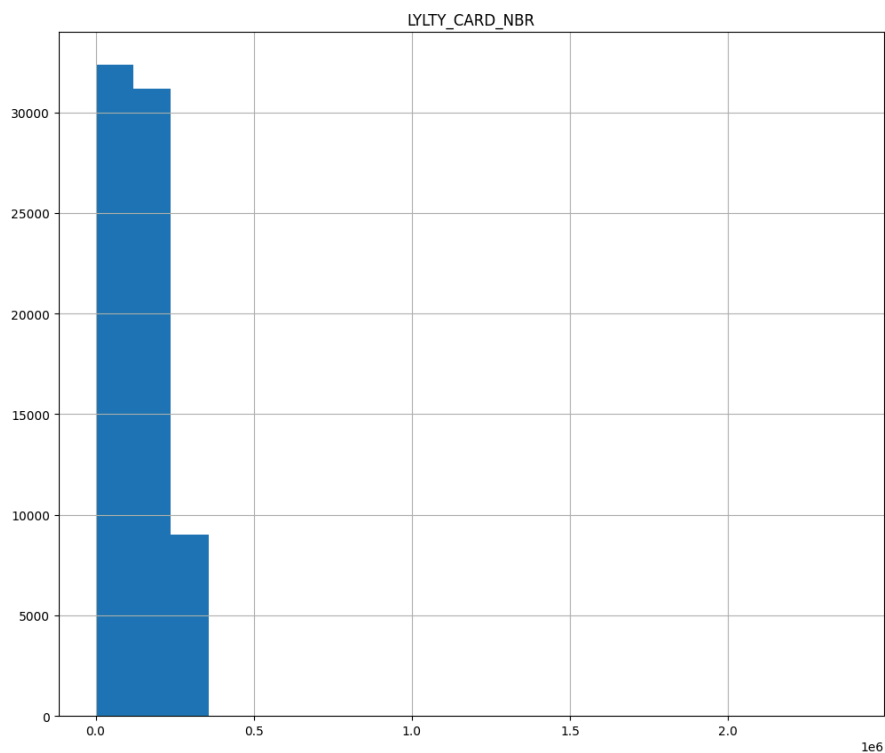
# Visualize outliers using boxplots
plt.figure(figsize=(12, 6))
sns.boxplot(data=df2)
plt.title('Boxplot of each column to check for outliers')
plt.xticks(rotation=45)
plt.show()
```



```
# Visualize outliers using histograms
plt.figure(figsize=(12, 6))
df2.hist(bins=20, figsize=(12, 10))
plt.suptitle('Histograms of each column to check for outliers')
plt.show()
```


<Figure size 1200x600 with 0 Axes>

Histograms of each column to check for outliers



```
# Visualize outliers using scatter plots
plt.figure(figsize=(12, 8))
for i, column in enumerate(df2.columns):
    plt.subplot(2, 4, i+1)
    plt.scatter(df2.index, df2[column])
    plt.title(column)
plt.suptitle('Scatter plots of each column to check for outliers')
plt.tight_layout()
plt.show()
```





```
# Removing duplicates
# Drop duplicate rows if any
df2.drop_duplicates(inplace=True)
```

Merging the two dataset

```
# Merge the two datasets based on a common column (in this case, 'LYLTY_CARD_NBR')
merged_df = pd.merge(df_cleaned, df2, on='LYLTY_CARD_NBR', how='inner')
print(merged_df)
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	\
0	1970-01-01 00:00:00.000043390	1	1000	1	
1	1970-01-01 00:00:00.000043599	1	1307	348	
2	1970-01-01 00:00:00.000043414	1	1307	346	
3	1970-01-01 00:00:00.000043533	1	1307	347	
4	1970-01-01 00:00:00.000043605	1	1343	383	
...	...	...	...	...	...
264830	1970-01-01 00:00:00.000043533	272	272319	270088	
264831	1970-01-01 00:00:00.000043325	272	272358	270154	
264832	1970-01-01 00:00:00.000043410	272	272379	270187	
264833	1970-01-01 00:00:00.000043461	272	272379	270188	
264834	1970-01-01 00:00:00.000043365	272	272380	270189	

	PROD_NBR	PROD_NAME	PROD_QTY	\
0	5	Natural Chip Compny SeaSalt175g	2	
1	66	CCs Nacho Cheese 175g	3	
2	96	WW Original Stacked Chips 160g	2	
3	54	CCs Original 175g	1	
4	61	Smiths Crinkle Cut Chips Chicken 170g	2	
...	...	...	...	...
264830	89	Kettle Sweet Chilli And Sour Cream 175g	2	
264831	74	Tostitos Splash Of Lime 175g	1	
264832	51	Doritos Mexicana 170g	2	
264833	42	Doritos Corn Chip Mexican Jalapeno 150g	2	
264834	74	Tostitos Splash Of Lime 175g	2	

	TOT_SALES	LIFESTAGE	PREMIUM_CUSTOMER
0	6.0	YOUNG SINGLES/COUPLES	Premium
1	6.3	MIDAGE SINGLES/COUPLES	Budget
2	3.8	MIDAGE SINGLES/COUPLES	Budget
3	2.1	MIDAGE SINGLES/COUPLES	Budget
4	2.9	MIDAGE SINGLES/COUPLES	Budget
...	...	...	...
264830	10.8	YOUNG SINGLES/COUPLES	Premium
264831	4.4	YOUNG SINGLES/COUPLES	Premium
264832	8.8	YOUNG SINGLES/COUPLES	Premium
264833	7.8	YOUNG SINGLES/COUPLES	Premium
264834	8.8	YOUNG SINGLES/COUPLES	Premium

[264835 rows x 10 columns]

Data Analysis

```

# Total Sales Analysis
total_sales = merged_df['TOT_SALES'].sum()
print("Total Sales:", total_sales)

Total Sales: 1934409.0

# Drivers of Sales
sales_by_product = merged_df.groupby('PROD_NAME')['TOT_SALES'].sum().sort_values(ascending=False).head(10)
print("Top 10 Products by Sales:")
print(sales_by_product)

Top 10 Products by Sales:
PROD_NAME
Dorito Corn Chp      Supreme 380g      40352.0
Smiths Crinkle Chip  Orgnl Big Bag 380g  36367.6
Smiths Crinkle Chips Salt & Vinegar 330g  34804.2
Kettle Mozzarella    Basil & Pesto 175g  34457.4
Smiths Crinkle        Original 330g      34302.6
Cheezels Cheese 330g      34296.9
Doritos Cheese        Supreme 330g      33390.6
Kettle Sweet Chilli And Sour Cream 175g  33031.8
Kettle Original 175g      32740.2
Kettle Sea Salt        And Vinegar 175g  32589.0
Name: TOT_SALES, dtype: float64

sales_by_store = merged_df.groupby('STORE_NBR')['TOT_SALES'].sum().sort_values(ascending=False).head(10)
print("Top 10 stores by sales:")
print(sales_by_store)

Top 10 stores by sales:
STORE_NBR
226      18905.45
88       16333.25
165      15973.75
40       15559.50
237      15539.50
58       15251.45
199      14797.00
4        14647.65
203      14551.60
26       14469.30
Name: TOT_SALES, dtype: float64

# Create a cross-tabulation of LIFESTAGE and PREMIUM_CUSTOMER
crosstab = pd.crosstab(merged_df['LIFESTAGE'], merged_df['PREMIUM_CUSTOMER'])
print("Crosstabulation of LIFESTAGE and PREMIUM_CUSTOMER:")
print(crosstab)

# Calculate the percentage of premium customers within each LIFESTAGE
crosstab_percentage = crosstab.div(crosstab.sum(axis=1), axis=0)
print("\nPercentage of premium customers within each LIFESTAGE:")
print(crosstab_percentage)

# Segment customers based on LIFESTAGE and PREMIUM_CUSTOMER
merged_df['Customer_Segment'] = merged_df['LIFESTAGE'] + '_' + merged_df['PREMIUM_CUSTOMER'].astype(str)
print("\nCustomer Segmentation:")
print(merged_df['Customer_Segment'].value_counts())

Crosstabulation of LIFESTAGE and PREMIUM_CUSTOMER:
PREMIUM_CUSTOMER      Budget  Mainstream  Premium
LIFESTAGE
MIDAGE SINGLES/COUPLES    5020      11874      8216
NEW FAMILIES              3005       2325      1589
OLDER FAMILIES           23160      14244      11192
OLDER SINGLES/COUPLES    18407      18318      17753
RETIRES                   15201      21466      13096
YOUNG FAMILIES           19122      12907      11563
YOUNG SINGLES/COUPLES    9242       20854       6281

Percentage of premium customers within each LIFESTAGE:
PREMIUM_CUSTOMER      Budget  Mainstream  Premium
LIFESTAGE
MIDAGE SINGLES/COUPLES  0.199920    0.472879    0.327200
NEW FAMILIES           0.434311    0.336031    0.229657
OLDER FAMILIES         0.476582    0.293111    0.230307
OLDER SINGLES/COUPLES  0.337880    0.336246    0.325875
RETIRES                0.305468    0.431365    0.263167
YOUNG FAMILIES         0.438658    0.296086    0.265255
YOUNG SINGLES/COUPLES  0.254062    0.573274    0.172664

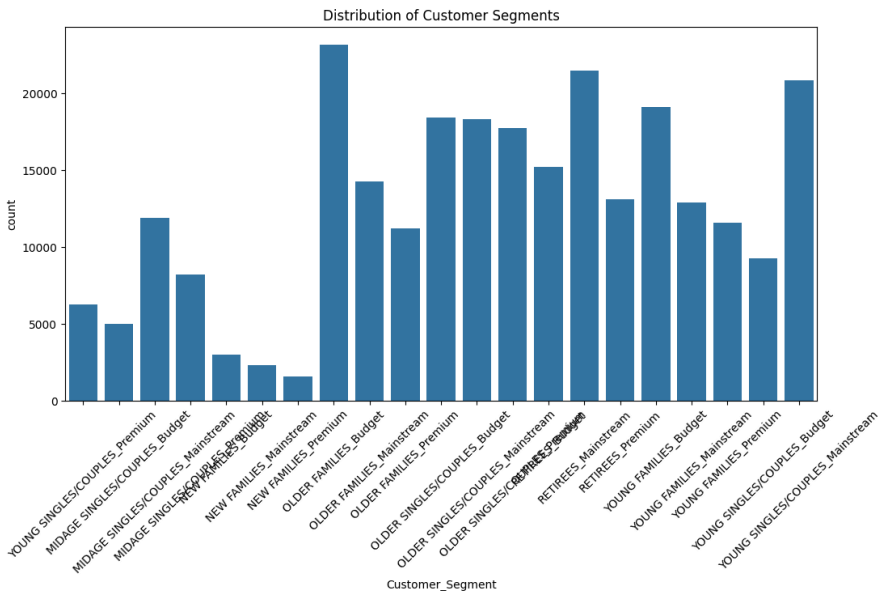
Customer Segmentation:
OLDER FAMILIES_Budget      23160
RETIRES_Mainstream        21466

```

YOUNG SINGLES/COUPLES_Mainstream	20854
YOUNG FAMILIES_Budget	19122
OLDER SINGLES/COUPLES_Budget	18407
OLDER SINGLES/COUPLES_Mainstream	18318
OLDER SINGLES/COUPLES_Premium	17753
RETIREEES_Budget	15201
OLDER FAMILIES_Mainstream	14244
RETIREEES_Premium	13096
YOUNG FAMILIES_Mainstream	12907
MIDAGE SINGLES/COUPLES_Mainstream	11874
YOUNG FAMILIES_Premium	11563
OLDER FAMILIES_Premium	11192
YOUNG SINGLES/COUPLES_Budget	9242
MIDAGE SINGLES/COUPLES_Premium	8216
YOUNG SINGLES/COUPLES_Premium	6281
MIDAGE SINGLES/COUPLES_Budget	5020
NEW FAMILIES_Budget	3005
NEW FAMILIES_Mainstream	2325
NEW FAMILIES_Premium	1589

Name: Customer\_Segment, dtype: int64

```
# Plot the distribution of customer segments
plt.figure(figsize=(12, 6))
sns.countplot(data=merged_df, x='Customer_Segment')
plt.title('Distribution of Customer Segments')
plt.xticks(rotation=45)
plt.show()
```



- Young and Non-Premium:**
  - This segment is likely to be price-sensitive and looking for value.
  - They may be more likely to purchase private label brands or generic products.
  - Consider offering discounts, promotions, and loyalty programs to attract and retain these customers.
- Young and Premium:**
  - This segment is likely to be more affluent and willing to pay for quality and convenience.
  - They may be more likely to purchase name-brand products and shop at higher-end stores.
  - Consider offering personalized shopping experiences, exclusive products, and premium services to cater to this segment.
- Established and Non-Premium:**

- This segment is likely to be more price-conscious and looking for value.
- They may be more likely to purchase private label brands or generic products.
- Consider offering discounts, promotions, and loyalty programs to attract and retain these customers.
- **Established and Premium:**
  - This segment is likely to be more affluent and willing to pay for quality and convenience.
  - They may be more likely to purchase name-brand products and shop at higher-end stores.
  - Consider offering personalized shopping experiences, exclusive products, and premium services to cater to this segment.
- **Seniors and Non-Premium:**
  - This segment is likely to be more price-sensitive and looking for value.
  - They may be more likely to purchase private label brands or generic products.
  - Consider offering discounts, promotions, and loyalty programs to attract and retain these customers.
- **Seniors and Premium:**
  - This segment is likely to be more affluent and willing to pay for quality and convenience.
  - They may be more likely to purchase name-brand products and shop at higher-end stores.
  - Consider offering personalized shopping experiences, exclusive products, and premium services to cater to this segment.

### Summary of Findings:

- The dataset contains transactional data for a retail store, including information about products, sales, and customer demographics.
- There were missing values and outliers in the dataset, which were handled through data cleaning techniques such as imputation and outlier removal.
- The total sales for the period were calculated.
- The top 10 products and stores by sales were identified.
- A cross-tabulation of LIFESTAGE and PREMIUM\_CUSTOMER was created to understand the relationship between these variables.
- Customers were segmented based on LIFESTAGE and PREMIUM\_CUSTOMER to identify different customer profiles.

### Overall Conclusion:

The analysis provides valuable insights into the sales performance, customer demographics, and customer segmentation for the retail store. The findings suggest that:

- The store should focus on promoting its top-selling products and stores to maximize sales.
- Different customer segments have distinct preferences and behaviors, which should be considered when developing marketing and promotional strategies.
- The store should consider offering personalized shopping experiences, exclusive products, and premium services to cater to the needs of premium customers.
- The store should implement loyalty programs and promotions to attract and retain price-sensitive customers.

By leveraging these insights, the retail store can optimize its product offerings, marketing strategies, and customer engagement efforts to improve sales and profitability.

```
# Create a line chart of total sales over time
plt.figure(figsize=(12, 6))
sns.lineplot(data=merged_df, x='DATE', y='TOT_SALES')
plt.title('Total Sales Over Time')
plt.xticks(rotation=45)
plt.show()
```



```
# Group the data by date and calculate the total sales and number of transactions for each day
daily_sales = merged_df.groupby('DATE')[['TOT_SALES', 'PROD_QTY']].sum().reset_index()
```

```
# Create a line chart of total sales over time
```